

# DATABASE SYSTEM PRINCIPLE - SQL QUERY LANGUAGE

李旭东 Li-Xudong

LEEUDONG@NANKAI.EDU.CN  
NANKAI UNIVERSITY

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## OVERVIEW OF THE SQL QUERY LANGUAGE

- IBM Sequel language developed as part of **System R** project at the IBM San Jose Research Laboratory
  - 1970s
  - Renamed Structured Query Language (SQL)
- In 1986, the American National Standards Institute (ANSI) and the International Organization for Standardization (ISO) published an SQL standard, called SQL-86

©LXD

## OVERVIEW OF THE SQL QUERY LANGUAGE

- Structured Query Language (SQL)
  - Data-Definition Language, DDL
  - Data-Manipulation Language, DML
  - Integrity 完整性
  - View Definition 视图定义
  - Transaction Control 事务控制
  - Embedded SQL and dynamic SQL 嵌入式和动态SQL
  - Authorization 授权
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.

©LXD

## ANSI AND ISO STANDARD SQL

- SQL-86 (or SQL-87) is the ISO 9075:1987 standard of 1987
- SQL-89 is the ISO/IEC 9075:1989 standard of 1989
- SQL-92 is the ISO/IEC 9075:1992 standard of 1992
- SQL:1999 is the ISO/IEC 9075:1999 standard of 1999
- SQL:2003 is the ISO/IEC 9075:2003 standard of 2003
- SQL:2006 is the ISO/IEC 9075:2006 standard of 2006
- SQL:2008 is the ISO/IEC 9075:2008 standard of 2008
- SQL:2011 is the ISO/IEC 9075:2011 standard of 2011

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- **Data Definition**
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## SQL DATA DEFINITION

- SQL DDL
  - The schema for each relation
  - The types of values associated with each attribute
  - The integrity constraints
  - The set of indices to be maintained for each relation
  - The security and authorization information for each relation
  - The physical storage structure of each relation on disk

©LXD

## DOMAIN TYPES IN SQL

- `char(n)`. Fixed length character string, with user-specified length  $n$ .
- `varchar(n)`. Variable length character strings, with user-specified maximum length  $n$ .
- `int`. Integer (a finite subset of the integers that is machine-dependent).
- `smallint`. Small integer (a machine-dependent subset of the integer domain type).
- `numeric(p,d)`. Fixed point number, with user-specified precision of  $p$  digits, with  $d$  digits to the right of decimal point. (ex., `numeric(3,1)`, allows 44.5, not 444.5 or 0.32)
- `real`, `double precision`. Floating point and double-precision floating point numbers, with machine-dependent precision.
- `float(n)`. Floating point number, with user-specified precision of at least  $n$  digits.

©LXD

## MYSQL DATA TYPES

- <https://dev.mysql.com/doc/refman/5.7/en/data-types.html>
  - Numeric Types
  - Date and Time Types
  - String Types
  - Spatial Data Types
  - The JSON Data Type
  - Data Type Default Values

©LXD

## BASIC SCHEMA DEFINITION

- Create table
- Insert into
- Delete from
- Alter table
- Drop table

©LXD

## CREATE TABLE CONSTRUCT

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1,
               A2 D2, ...,
               An Dn,
               (integrity-constraint1,完整性约束),
               ...,
               (integrity-constraintk))
```

- $r$  is the name of the relation
- each  $A_i$  is an attribute name in the schema of relation  $r$
- $D_i$  is the data type of values in the domain of attribute  $A_i$

©LXD

## CREATE TABLE CONSTRUCT

- the **create table** command
- Example:

```
create table instructor (
  ID      char(5),
  name    varchar(20),
  dept_name varchar(20),
  salary  numeric(8,2));
```

©LXD

## INTEGRITY CONSTRAINTS IN CREATE TABLE

- **not null**
- **primary key** ( $A_1, \dots, A_n$ )
- **foreign key** ( $A_m, \dots, A_n$ ) **references**  $r$

Example:

```
create table instructor (
  ID          char(5),
  name        varchar(20) not null,
  dept_name   varchar(20),
  salary      numeric(8,2),
  primary key (ID),
  foreign key (dept_name) references department);
```

primary key declaration on an attribute automatically ensures not null

©LXD

## CASES: CREATE TABLE

```
create table department
(dept_name varchar (20),
building varchar (15),
budget numeric (12,2),
primary key (dept name);
```

```
create table course
(course_id varchar (7),
title varchar (50),
dept_name varchar (20),
credits numeric (2,0),
primary key (course id),
foreign key (dept_name) references department);
```

©LXD

## CASES: CREATE TABLE

```
create table section
(course_id varchar (8),
sec_id varchar (8),
semester varchar (6),
year numeric (4,0),
building varchar (15),
room_number varchar (7),
time_slot_id varchar (4),
primary key (course_id, sec_id, semester, year),
foreign key (course_id) references course);
```

©LXD

## CASES: CREATE TABLE

```
create table teaches
(ID varchar (5),
course_id varchar (8),
sec_id varchar (8),
semester varchar (6),
year numeric (4,0),
primary key (ID, course_id, sec_id, semester, year),
foreign key (course_id, sec_id, semester, year) references section,
foreign key (ID) references instructor);
```

©LXD

## UPDATES TO TABLES

- **Insert**
  - **insert into** *instructor* values ('10211', 'Smith', 'Biology', 66000);
- **Delete**
  - Remove all tuples from the *student* relation
    - **delete from** *student*
- **Drop Table**
  - **drop table**  $r$

©LXD

## UPDATES TO TABLES

- **Alter**
  - **alter table**  $r$  **add**  $A$   $D$ 
    - where  $A$  is the name of the attribute to be added to relation  $r$  and  $D$  is the domain of  $A$ .
    - All exiting tuples in the relation are assigned *null* as the value for the new attribute.
  - **alter table**  $r$  **drop**  $A$ 
    - where  $A$  is the name of an attribute of relation  $r$
    - Dropping of attributes not supported by many databases.

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- **Basic Query Structure**
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## BASIC STRUCTURE OF SQL QUERIES

- A typical SQL query has the form:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

$A_i$  represents an attribute

- $R_i$  represents a relation
- $P$  is a predicate.
- The result of an SQL query is a **relation**.

©LXD

## QUERIES ON A SINGLE RELATION

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:
 

```
select name
from instructor
```
- NOTE: SQL names are case insensitive
  - E.g.,  $Name \equiv NAME \equiv name$
  - Some people use upper case wherever we use bold font.

©LXD

name
Srinivasan
Wu
Mozart
Einstein
El Said
Gold
Katz
Califieri
Singh
Crick
Brandt
Kim

## QUERIES ON A SINGLE RELATION

### - DEDUPLICATION 去重

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates
 

```
select distinct dept_name
from instructor
```
- The keyword **all** specifies that duplicates should not be removed.
 

```
select all dept_name
from instructor
```

©LXD

## QUERIES ON A SINGLE RELATION

### - MORE

- An **asterisk (\*)** in the select clause denotes "all attributes"
 

```
select *
from instructor
```
- An attribute can be a **literal** with no **from** clause
 

```
select '437'
```

  - Results is a table with one column and a single row with value "437"
  - Can give the column a name using:
 

```
select '437' as FOO
```

©LXD

## QUERIES ON A SINGLE RELATION

### - MORE

- An attribute can be a literal with **from** clause
 

```
select 'A'
from instructor
```

  - Result is a table with **one column** and  **$N$  rows** (number of tuples in the *instructors* table), each row with value "A"

©LXD

## QUERIES ON A SINGLE RELATION

### - MORE

- The **select** clause can contain arithmetic expressions involving the operation, +, -, \*, and /, and operating on constants or attributes of tuples.

- The query:

```
select ID, name, salary/12
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename "*salary/12*" using the *as* clause:

```
select ID, name, salary/12 as monthly_salary
```

©LXD

## THE WHERE CLAUSE

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```

©LXD

## THE WHERE CLAUSE

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000
 

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary > 80000
```
- Comparisons can be applied to results of arithmetic expressions.

©LXD

## QUERIES ON MULTIPLE RELATIONS

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P;
```

- Each  $A_i$  represents an attribute, and each  $r_i$  a relation.  $P$  is a predicate.
  - If the where clause is omitted, the predicate  $P$  is true.

©LXD

## QUERIES ON MULTIPLE RELATIONS

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor X teaches*

```
select *
from instructor, teaches
```

- generates every possible instructor – teaches pair, with all attributes from both relations.
- For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)

©LXD

## CARTESIAN PRODUCT 笛卡尔积

ID	name	dept_name	salary	ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	MU-199	1	Spring	2010
				22222	PHY-101	1	Fall	2009

inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...

©LXD

## CASES: CARTESIAN PRODUCT

- Find the names of all instructors who have taught some course and the course\_id
  - `select name, course_id`  
`from instructor, teaches`  
`where instructor.ID = teaches.ID`
- Find the names of all instructors in the Art department who have taught some course and the course\_id
  - `select name, course_id`  
`from instructor, teaches`  
`where instructor.ID = teaches.ID and instructor.dept_name = 'Art'`

©LXD

## THE NATURAL JOIN

```
select A1, A2, ..., An
  from r1 natural join r2 natural join ... natural join rm
 where P;
```

More generally, a from clause can be of the form

```
from E1, E2, ..., En
```

- where each E<sub>i</sub> can be a single relation or an expression involving natural joins.

©LXD

## CASES: THE NATURAL JOIN

The natural join of the instructor relation with the teaches relation

ID	name	dept_name	salary	ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	Wu	Finance	90000	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	Mozart	Music	40000	MU-199	1	Spring	2010
				22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
				32343	El Said	History	60000	HIS-351	1	Spring	2010
				45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
				45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
				76766	Crick	Biology	72000	BIO-101	1	Summer	2009
				76766	Crick	Biology	72000	BIO-301	1	Summer	2010
				83821	Brandt	Comp. Sci.	92000	CS-190	1	Spring	2009
				83821	Brandt	Comp. Sci.	92000	CS-190	2	Spring	2009
				83821	Brandt	Comp. Sci.	92000	CS-319	2	Spring	2010
				98345	Kim	Elec. Eng.	80000	EE-181	1	Spring	2009

©LXD

## CASES: THE NATURAL JOIN

- Way 1
 

```
select name, course_id
  from instructor, teaches
 where instructor.ID = teaches.ID;
```
- Way 2
 

```
select name, course_id
  from instructor natural join teaches;
```

©LXD

## CASES 2: THE NATURAL JOIN

- `select name, title`  
`from instructor natural join teaches, course`  
`where teaches.course_id = course.course_id;`
- `select name, title`  
`from instructor natural join teaches natural join course;`
- `select name, title`  
`from (instructor natural join teaches) join course using (course_id);`

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD



## THE RENAME OPERATION

- The SQL allows renaming relations and attributes using the **as** clause:  
*old-name as new-name*
- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.
  - ```
select distinct T.name
  from instructor as T, instructor as S
 where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```
- Keyword **as** is optional and may be omitted  
*instructor as T ≡ instructor T*

©LXD

## STRING OPERATIONS

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.

©LXD

## STRING OPERATIONS

- Find the names of all instructors whose name includes the substring "dar".
 

```
select name
  from instructor
 where name like '%dar%'
```
- Match the string "100%"
 

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.

©LXD

## STRING OPERATIONS

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro".
  - '%Comp%' matches any string containing "Comp" as a substring.
  - '\_\_\_' matches any string of exactly three characters.
  - '\_\_\_%' matches any string of at least three characters.

©LXD

## STRING OPERATIONS

- SQL supports a variety of string operations such as
  - concatenation (using "||")
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

©LXD

## ORDERING THE DISPLAY OF TUPLES

- List in alphabetic order the names of all instructors
 

```
select distinct name
  from instructor
 order by name
```
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - Example: **order by name desc**
- Can sort on multiple attributes
  - Example: **order by dept\_name, name**

©LXD

## WHERE CLAUSE PREDICATES

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )
  - `select name`  
`from instructor`  
`where salary between 90000 and 100000`

©LXD

## WHERE CLAUSE PREDICATES

- Way 1:
  - `select name, course_id`  
`from instructor, teaches`  
`where (instructor.ID= teaches.ID and dept_name = 'Biology');`
- Way 2:
  - `select name, course_id`  
`from instructor, teaches`  
`where (instructor.ID, dept_name) = (teaches.ID, 'Biology');`

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- **Set Operations**
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## SET OPERATIONS

- union
- intersect
- except

©LXD

## SET OPERATIONS

### - UNION

- Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
union
(select course_id from section where sem = 'Spring' and
year = 2010)
```

©LXD

## SET OPERATIONS

### - INTERSECT

Find courses that ran in Fall 2009 and in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
intersect
(select course_id from section where sem = 'Spring'
and year = 2010)
```

©LXD



## SET OPERATIONS

### - EXCEPT

Find courses that ran in Fall 2009 but not in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
except
(select course_id from section where sem = 'Spring'
and year = 2010)
```

©LXD

## SET OPERATIONS

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  $r$  **union all**  $s$
  - $\min(m, n)$  times in  $r$  **intersect all**  $s$
  - $\max(0, m - n)$  times in  $r$  **except all**  $s$

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values**
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## NULL VALUES

- It is possible for tuples to have a null value, denoted by **null**, for some of their attributes
- null** signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving **null** is **null**
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```

©LXD

## NULL VALUES AND THREE VALUED LOGIC

- Three values – **true**, **false**, **unknown**
- Any comparison with **null** returns **unknown**
  - Example:  $5 < \text{null}$  or  $\text{null} <= \text{null}$  or  $\text{null} = \text{null}$

©LXD

## NULL VALUES AND THREE VALUED LOGIC

- Three-valued logic using the value **unknown**:
  - OR:  $(\text{unknown or true}) = \text{true}$ ,  
 $(\text{unknown or false}) = \text{unknown}$   
 $(\text{unknown or unknown}) = \text{unknown}$
  - AND:  $(\text{true and unknown}) = \text{unknown}$ ,  
 $(\text{false and unknown}) = \text{false}$ ,  
 $(\text{unknown and unknown}) = \text{unknown}$
  - NOT:  $(\text{not unknown}) = \text{unknown}$
  - " $P$  is **unknown**" evaluates to true if predicate  $P$  evaluates to **unknown**
- Result of **where** clause predicate is treated as **false** if it evaluates to **unknown**

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- **Aggregate Functions**
- Nested Subqueries
- Modification of the Database

©LXD

## AGGREGATE FUNCTIONS 聚集函数

- These functions operate on the multiset of values of a column of a relation, and return a value
  - avg:** average value
  - min:** minimum value
  - max:** maximum value
  - sum:** sum of values
  - count:** number of values

©LXD

## AGGREGATE FUNCTIONS

- Find the average salary of instructors in the Computer Science department
  - `select avg (salary)`  
from `instructor`  
where `dept_name='Comp. Sci.'`;
- Find the total number of instructors who teach a course in the Spring 2010 semester
  - `select count (distinct ID)`  
from `teaches`  
where `semester = 'Spring' and year = 2010;`

©LXD

## AGGREGATE FUNCTIONS – GROUP BY

- Find the average salary of instructors in each department
  - `select dept_name, avg (salary) as avg_salary`  
from `instructor`  
group by `dept_name;`

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 38583 | Calitieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

| dept_name  | salary |
|------------|--------|
| Biology    | 72000  |
| Comp. Sci. | 77333  |
| Elec. Eng. | 80000  |
| Finance    | 85000  |
| History    | 61000  |
| Music      | 40000  |
| Physics    | 91000  |

©LXD

## AGGREGATE FUNCTIONS – GROUP BY

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
  - **/\* erroneous query \*/**  
`select dept_name, ID, avg (salary)`  
from `instructor`  
group by `dept_name;`

©LXD

## AGGREGATE FUNCTIONS – HAVING CLAUSE

- Find the names and average salaries of all departments whose average salary is greater than 42000
  - `select dept_name, avg (salary)`  
from `instructor`  
group by `dept_name`  
having `avg (salary) > 42000;`

Note: predicates in the having clause are applied after the formation of groups whereas predicates in the where clause are applied before forming groups

©LXD

## AGGREGATE FUNCTIONS

- Total all salaries

```
select sum (salary )
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount

©LXD

## AGGREGATE FUNCTIONS

- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- **Nested Subqueries**
- Modification of the Database

©LXD

## NESTED SUBQUERIES 嵌套子查询

- SQL provides a mechanism for nesting subqueries.
- A subquery is a select-from-where expression that is nested within another query.

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
      course_id in (select course_id
                    from section
                    where semester = 'Spring' and year= 2010);
```

©LXD

## NESTED SUBQUERIES 嵌套子查询

- The nesting can be done in the following SQL query
- ```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

as follows:

- A<sub>i</sub> can be replaced by a subquery that generates a single value.
- r<sub>i</sub> can be replaced by any valid subquery
- P can be replaced with an expression of the form:

B <operation> (subquery)

Where B is an attribute and <operation> to be defined later.

©LXD

## SUBQUERIES IN THE WHERE CLAUSE

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality.

©LXD

## NESTED SUBQUERIES

## - SET MEMBERSHIP 集合成员测试

## • in, not in

Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id not in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

©LXD

## NESTED SUBQUERIES

## - SET MEMBERSHIP

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
select count (distinct ID)
from takes
where (course_id, sec_id, semester, year) in
(select course_id, sec_id, semester, year
from teaches
where teaches.ID= 10101);
```

©LXD

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept name = 'Biology';
```

Same query using > some clause

```
select name
from instructor
where salary > some (select salary
from instructor
where dept name = 'Biology');
```

©LXD

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

- $F < \text{comp} > \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F < \text{comp} > t)$

Where <comp> can be: <, ≤, >, =, ≠

- > some
- < some
- <= some
- >= some
- <> some

◆ (not same to not in)

- = some (same to in)

©LXD

$(5 < \text{some } \begin{smallmatrix} 0 \\ 5 \\ 6 \end{smallmatrix})$	$= \text{true}$	(read: 5 < some tuple in the relation)
$(5 < \text{some } \begin{smallmatrix} 0 \\ 5 \end{smallmatrix})$	$= \text{false}$	
$(5 = \text{some } \begin{smallmatrix} 0 \\ 5 \end{smallmatrix})$	$= \text{true}$	
$(5 \neq \text{some } \begin{smallmatrix} 0 \\ 5 \end{smallmatrix})$	$= \text{true (since } 0 \neq 5)$	
$(= \text{some})$	$= \text{in}$	However, $(\neq \text{some}) \neq \text{not in}$

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

- > all
- < all
- <= all
- >= all
- = all (not same to in)
- <> all (same to not in)

- $F < \text{comp} > \text{all } r \Leftrightarrow \forall t \in r (F < \text{comp} > t)$

(5 < all	<table><tr><td>0</td></tr><tr><td>5</td></tr><tr><td>6</td></tr></table> )	0	5	6	= false
0					
5					
6					
(5 < all	<table><tr><td>6</td></tr><tr><td>10</td></tr></table> )	6	10	= true	
6					
10					
(5 = all	<table><tr><td>4</td></tr><tr><td>5</td></tr></table> )	4	5	= false	
4					
5					
(5 ≠ all	<table><tr><td>4</td></tr><tr><td>6</td></tr></table> )	4	6	= true (since 5 ≠ 4 and 5 ≠ 6)	
4					
6					
(≠ all) = not in					
However, (= all) ≠ in					

©LXD

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

?What meaning?

```
select dept_name
from instructor
group by dept_name
having avg (salary) >= all (select avg (salary)
from instructor
group by dept_name);
```

©LXD

## TEST FOR EMPTY RELATIONS

- exists , not exists
- Yet another way of specifying the query "Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester"
 

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
exists (select *
from section as T
where semester = 'Spring' and year = 2010 and S.course_id = T.course_id);
```
- **Correlation name** – variable S in the outer query
- **Correlated subquery** – the inner query

©LXD

## TEST FOR EMPTY RELATIONS

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name from student as S
where not exists ((select course_id from course
where dept_name = 'Biology')
except
(select T.course_id from takes as T
where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took

Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

Note: Cannot write this query using = all and its variants

©LXD

## TEST FOR ABSENCE OF DUPLICATE TUPLES

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to "true" if a given subquery contains no duplicates.
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
from section as R
where T.course_id = R.course_id
and R.year = 2009);
```

©LXD

## SUBQUERIES IN THE FROM CLAUSE

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
from instructor
group by dept_name)
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause

©LXD

## SUBQUERIES IN THE FROM CLAUSE

- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."
- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
from instructor
group by dept_name) as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

©LXD

## WITH CLAUSE

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs
  - Introduced in SQL:1999
- Find all departments with the maximum budget
 

```
with max_budget (value) as
(select max(budget)
from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

©LXD

## COMPLEX QUERIES USING WITH CLAUSE

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total(dept_name, value) as
(select dept_name, sum(salary)
 from instructor
 group by dept_name),
dept_total_avg(value) as
(select avg(value)
 from dept_total)
select dept_name
 from dept_total, dept_total_avg
 where dept_total.value > dept_total_avg.value;
```

©LXD

## SCALAR SUBQUERIES 标量子查询

- Scalar subqueries
  - SQL allows subqueries to occur wherever an expression returning a value is permitted, provided the subquery returns only one tuple containing a single attribute
  - List all departments along with the number of instructors in each department

```
select dept_name,
       (select count(*)
        from instructor where department.dept_name = instructor.dept_name)
       as num_instructors
  from department;
```

- Runtime error if subquery returns more than one result tuple

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## MODIFICATION OF THE DATABASE

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

©LXD

## DELETION 删除

- Delete all instructors
 

```
delete from instructor
```
- Delete all instructors from the Finance department
 

```
delete from instructor
 where dept_name = 'Finance';
```
- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.
 

```
delete from instructor
 where dept_name in (select dept_name
                    from department
                    where building = 'Watson');
```

©LXD

## DELETION

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor
 where salary < (select avg(salary)
                from instructor);
```

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  - First, compute avg(salary) and find all tuples to delete
  - Next, delete all tuples found above (without recomputing avg or retesting the tuples)

©LXD



## INSERTION插入

- Add a new tuple to *course*

```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student
values ('3003', 'Green', 'Finance', null);
```

©LXD

## INSERTION

- Add all instructors to the *student* relation with *tot\_creds* set to 0

```
insert into student
select ID, name, dept_name, 0
from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
would cause problem
```

©LXD

## UPDATES更新

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%

- Write two **update** statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
```

```
update instructor
set salary = salary * 1.05
where salary <= 100000;
```

- *The order is important*

- Can be done better using the **case** statement (next slide)

©LXD

## UPDATE

### - CASE STATEMENT FOR CONDITIONAL UPDATES

- Same query as before but with **case** statement

```
update instructor
set salary = case
when salary <= 100000 then salary * 1.05
else salary * 1.03
end
```

©LXD

## UPDATE

### - CASE STATEMENT FOR CONDITIONAL UPDATES

- The general form of the case statement is as follows

```
case
when pred1 then result1
when pred2 then result2
...
when predn then resultn
else result0
end
```

©LXD

## UPDATES WITH SCALAR SUBQUERIES

- Recompute and update *tot\_creds* value for all students

```
update student S
set tot_cred = (select sum(credits)
from takes, course
where takes.course_id = course.course_id and
S.ID = takes.ID and takes.grade <> 'F' and
takes.grade is not null);
```

- Sets *tot\_creds* to null for students who have not taken any course

©LXD

## UPDATES WITH SCALAR SUBQUERIES

- Recompute and update tot\_creds value for all students
- Instead of `sum(credits)`, use:

```
update student S
set tot_cred =
  case
  when sum(credits) is not null then sum(credits)
  else 0
  end
```

©LXD

## SUMMARY

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

Q&A?

THANKS!

leexudong@nankai.edu.cn